Interview prep guide - Forward Deployed Architect, DevRev

Welcome! We're excited to have you in our interview process. This guide will help you understand what to expect at each stage and how to prepare.

About DevRev

DevRev is building the future of work with **Computer** - your AI teammate. We unify data sources, tools, and workflows into a single AI-ready platform, enabling human-AI collaboration that breaks down silos and eliminates repetitive work. Backed by Khosla Ventures and Mayfield with $150M+ raised, DevRev is trusted by global companies across industries.

About the role

As a Forward Deployed Architect, you'll be a hands-on senior technical architect on our Applied AI Engineering team. You'll own end-to-end design and delivery of AI-driven business transformation projects - from scoping with pre-sales teams to building proofs-of-concept, custom integrations, and production deployments. Expect to spend 20-40% of your time working directly with customers.

Interview process overview

Stage	Format	Duration
1. Recruiter screening	Google Meet (audio)	30-45 min
2. Hiring Manager screening	Video call	45-60 min

3. Coding round	Live coding	60 min
4. System design round	Whiteboard / virtual	60 min
5. Founder's mentality round	Conversation	45 min
6. You.io assessment	Online	Self-paced
7. Reference check	Async	-

Round 1: Recruiter screening

What to expect

A conversation about your background, career motivations, and interest in the role. We'll also share more about DevRev and answer your questions.

How to prepare

- **Know your story.** Be ready to walk through your career arc in 3-5 minutes - what's the thread connecting your moves?
- **Reflect on your technical breadth.** Where are you strongest - backend systems, AI/ML, cloud infra, or customer-facing solution design?
- **AI experience.** Be prepared to talk about your most hands-on experience with LLMs, RAG, or agentic AI systems.
- **Scale of impact.** Think about the largest system you've architected or the most complex integration you've delivered.
- **Coding recency.** We'll ask how hands-on you are today and what languages you're writing production code in.
- **Motivation.** Why this role? What excites you about being forward-deployed and customer-facing vs. a pure engineering role?
- **Research DevRev.** Explore our website, product, and recent news. Understand what Computer does.

Tips

- Be concise and specific - use concrete examples over general statements.
- It's OK to ask us questions too! This is a two-way conversation.
- Be transparent about your notice period, compensation expectations, and location preferences.

Round 2: Hiring Manager screening

What to expect

A deeper conversation focused on three areas: (1) how you work with external stakeholders, (2) team leadership and cross-functional collaboration, and (3) your agentic AI expertise.

How to prepare

External stakeholder management

- Prepare 2-3 stories about owning a technical relationship with a customer or partner.
- Think about a time a customer requirement was ambiguous or conflicting - how did you resolve it?
- Recall a situation where you had to push back on a customer's technical approach diplomatically.
- If you've run live demos or POCs for customers, be ready to walk through the experience.

Team leadership & cross-functional work

- Prepare an example of leading delivery across engineering, product, and customer success.
- Think about a time you aligned stakeholders with competing priorities.
- How do you handle limited engineering capacity against a fixed customer deadline?

Agentic AI depth

- Be ready to describe agentic AI systems you've built or deployed - walk through the architecture.
- How do you evaluate whether an AI agent is working correctly? What metrics and evals do you use?
- Think about how you'd architect a multi-agent system integrating with CRM, ticketing, and knowledge base.
- How do you approach prompt engineering at scale - versioning, testing, iteration?
- How do you handle hallucination and reliability concerns in enterprise deployments?

Tips

- Use the STAR format (Situation, Task, Action, Result) for behavioral stories.
- Show that you can be both a trusted advisor and a hands-on builder.
- Demonstrate production-grade thinking, not just prototyping.

Round 3: Coding round

What to expect

A live coding session (shared screen or CoderPad). You'll solve 1-2 problems in real time.

Languages

Python or TypeScript preferred. Go is a bonus.

What we assess

- Clean, readable, well-structured code.
- Edge case handling and error resilience.
- Communication while coding - explain your thinking as you go.
- Problem decomposition and clarifying questions before jumping in.

Areas to brush up on

- **API integration:** Calling external APIs, handling pagination, retries, rate limits.
- **Data transformation:** Parsing and transforming nested JSON payloads into normalized schemas.
- **Agent logic:** Implementing a tool-calling loop - given a user query, select the right tool, call it, synthesize a response.
- **Core DSA:** Graph traversal, tree-based problems, hash maps, queues - especially patterns relevant to workflow orchestration.

Tips

- Think out loud. We want to understand your approach, not just see the final answer.
- Ask clarifying questions before you start coding.
- Start with a working solution, then optimize - don't over-engineer from the start.
- Handle errors gracefully - this role involves production systems.

Round 4: System design round

What to expect

You'll be given an open-ended architecture problem and asked to design a solution on a whiteboard (or virtual equivalent). Problems will be AI-native and customer-facing in nature.

Example problem types

- Designing an AI agent platform for non-technical users to configure workflow automations across multiple channels.
- Architecting a RAG-based system that pulls from multiple enterprise data sources with access control.
- Designing a delivery framework that takes a customer from scoping doc to deployed AI agent in under 2 weeks.

What we assess

- Clarity of high-level architecture.
- Trade-off articulation (why this approach over alternatives).
- Integration and scalability thinking.
- Security, observability, and failure handling.
- Customer-centric design - not just engineering elegance.

How to prepare

- Practice designing systems that involve LLM orchestration, multi-agent coordination, and external API integrations.
- Be comfortable discussing RAG architectures, vector databases, embedding strategies, and eval pipelines.
- Think about multi-tenancy, data isolation, and access control in AI systems.
- Review patterns for CI/CD of AI agents (prompt versioning, A/B testing, rollback).
- Prepare to discuss monitoring and observability for LLM-based systems.

Tips

- Start with requirements clarification and scope definition.
- Draw clear component diagrams - label everything.
- Explicitly state trade-offs at each decision point.
- Think about the customer experience, not just the backend.
- Address failure modes and how the system degrades gracefully.

Round 5: Founder's mentality round

What to expect

A conversation about your ownership mindset, bias for action, resourcefulness, and alignment with how we work at DevRev.

How to prepare

- **0-to-1 building:** Prepare a story about building something from scratch with minimal guidance.
- **Constructive disagreement:** Think about a time you disagreed with leadership - what did you do?
- **Learning velocity:** When did you last learn a completely new technology or domain in under 2 weeks?
- **Urgency & preparation:** If your first customer meeting was in 3 days, how would you prepare?
- **Industry perspective:** What's broken in how companies deploy AI today? What would you do differently?
- **Failure & growth:** Be ready to talk about a failure - what you owned and what you changed.

What we're looking for

- Ownership over outcomes, not just tasks.
- Bias for action in ambiguous situations.
- Intellectual curiosity and speed of learning.
- Customer obsession.
- No-ego collaboration.

Tips

- Be authentic - we're not looking for rehearsed answers.
- Show that you take initiative without waiting for permission.
- Demonstrate that you care about impact, not just execution.

General tips for all rounds

1. **Be specific.** Vague answers don't land. Use real examples with concrete details.
2. **Show customer empathy.** This role is about solving problems for real people - demonstrate that mindset.

3. **Balance depth and breadth.** Show that you can go deep technically while also seeing the big picture.
4. **Ask great questions.** The best candidates are curious about our challenges and how they can contribute.
5. **Be yourself.** We're hiring a person, not a performance.

Key skills we value for this role

Must-have	Nice-to-have
Applied AI depth (LLMs, agents, RAG)	Enterprise architecture certs (TOGAF, SAFe)
Hands-on coding (Python/TypeScript)	Go proficiency
Customer-facing confidence	Startup experience
System design at scale	
Cross-functional leadership	
Ownership mindset	

Resources to explore before your interviews

- DevRev website - understand our product and vision.
- Familiarize yourself with agentic AI patterns - tool calling, function calling, multi-agent orchestration.
- Review RAG architectures and evaluation frameworks for LLM systems.
- Read about DevRev's Computer and how it differs from traditional SaaS tools.

Good luck with your preparation! We look forward to getting to know you.